

nmf

October 4, 2023

```
[5]: """
Inference & Representation - 2022 Fall HW1
Question 2 PCA and Non-negative matrix factorization.

"""

"""
Tools for loading the MNIST Data.
From Optimization Based Data Analysis HW1
@author: Brett
"""

import numpy as np
from mnist_tools import *
from plot_tools import *
import matplotlib.pyplot as plt

"""
Given train (in the format returned by load_train_data in mnist_tools),
and a 1d numpy array testImage you should return a tuple (digit,imageIdx).
    ↳digit is
an integer giving the numerical digit value of the training image closest
to the testImage in Euclidean distance. imageIdx is the row number of the
    ↳closest
training image in the 2d array train[digit].
"""


"""
Assumes the data file is in 'mnist_all.mat'.
"""

datafile = "mnist_all.mat" #Change if you put the file in a different path
train = load_train_data(datafile)

trainarr=np.asarray(train)
trainarr = np.reshape(trainarr, (trainarr.shape[0]*trainarr.shape[1],-1))
trainarr = trainarr.astype(float)
trainarr=trainarr-trainarr.mean(axis=0)
```

```
[6]: """
Plot of the singular vectors corresponding
to top 10 singular values of the data.
@author: Vlad
"""
U, s, V = np.linalg.svd(trainarr, full_matrices=True)
n=10
imgs = [V[i,:] for i in range(n)]
plot_image_grid(imgs,
                "Singular vectors corresponding to top 10 singular values of the_
↳data")
```

Singular vectors corresponding to top 10 singular values of the data



```
[7]: """
Plot of the results of the nearest neighbour test applied
to a principal component projection.
@author: Vlad
"""

def project(V, Images) :
    return np.dot(V.T, np.dot(V, Images))

def compute_nearest_neighbors(train, testImage, V) :
    train=[np.array(i, dtype=float) for i in train]
    testImage= np.array(testImage, dtype=float)
    digit=0
    imageIdx=0
    dist=np.linalg.norm (project (V, train[digit][imageIdx])-project (V,
↳testImage))
    for i in range(len(train)):
        for j in range (train[i].shape[0]):
```

```

        tempDist=np.linalg.norm (project(V,train[i][j])-project(V,
↪testImage))
        if tempDist<dist:
            digit=i
            imageIdx =j
            dist= tempDist
        return digit, imageIdx


n=8
U, s, V = np.linalg.svd(trainarr, full_matrices=False)
V=V[0:n,:]

test,testLabels = load_test_data(datafile)

imgs = []
TestLabels = []
for i in range(len(testLabels)) :
    trueDigit = testLabels[i]
    testImage = test[i]
    (nnDig,nnIdx) = compute_nearest_neighbors(train,testImage,V)
    imgs.extend( [testImage,train[nnDig][nnIdx,:]] )
    TestLabels.append(nnDig)

row_titles = ['Test','Nearest']
col_titles = ['%d vs. %d'%(i,j) for i,j in zip(testLabels,testLabels)]
plot_image_grid(imgs,
                "Image-NearestNeighbor",
                ↪
↪(28,28),len(testLabels),2,True,row_titles=row_titles,col_titles=col_titles)

```



```
[8]: import numpy as np
import matplotlib.pyplot as plt
import scipy.io

# Loading the MNIST data
datafile = "mnist_all.mat"
mnist_data = scipy.io.loadmat(datafile)

# Checking the keys in the loaded data
mnist_data.keys()
```

```
[8]: dict_keys(['__header__', '__version__', '__globals__', 'train0', 'test0',
'train1', 'test1', 'train2', 'test2', 'train3', 'test3', 'train4', 'test4',
'train5', 'test5', 'train6', 'test6', 'train7', 'test7', 'train8', 'test8',
'train9', 'test9'])
```

```
[9]: # Extracting training and testing data and labels
def extract_data_and_labels(mnist_data):
    data = []
    labels = []
    for i in range(10):
        train_data_key = f'train{i}'
        test_data_key = f'test{i}'
        train_data = mnist_data[train_data_key]
        test_data = mnist_data[test_data_key]
        combined_data = np.vstack((train_data, test_data))
        data.append(combined_data)
        labels.append(np.full((combined_data.shape[0],), i))
    return data, labels

# Getting the data and labels
data, labels = extract_data_and_labels(mnist_data)

# Checking the shape of extracted data for a specific digit
data[0].shape, labels[0].shape
```

```
[9]: ((6903, 784), (6903,))
```

```
[12]: from sklearn.decomposition import NMF
import plot_tools

# Function to apply NMF and plot the components
def apply_nmf_and_plot(data, n_components):
    # Combining all digit data
    all_data = np.vstack(data)
```

```

# Applying NMF
nmf = NMF(n_components=n_components, init='random', random_state=0)
W = nmf.fit_transform(all_data)
H = nmf.components_

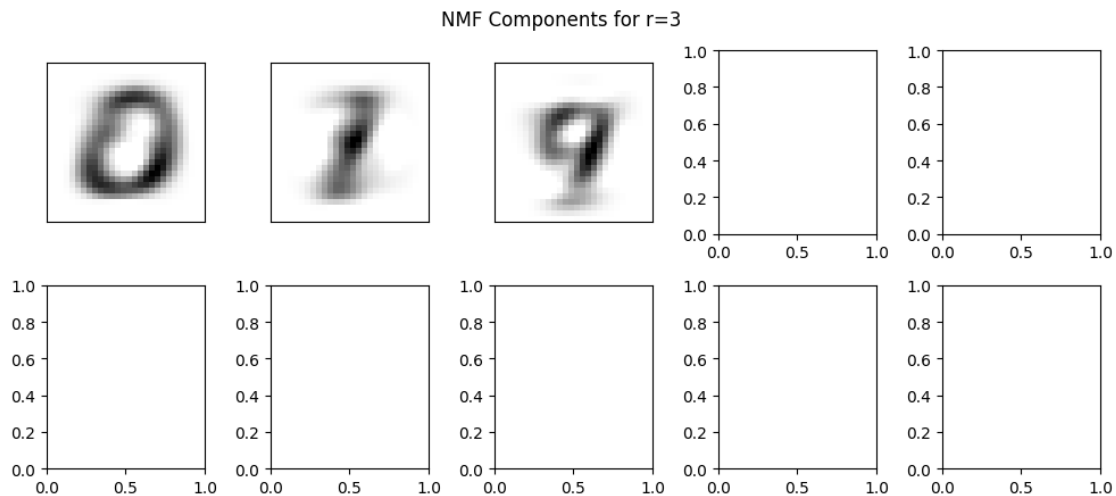
# Plotting the components
plot_tools.plot_image_grid(H, f"NMF Components for r={n_components}",
↪image_shape=(28, 28))

# Plotting function (from the provided plot_tools.py with minor modifications)
# def plot_image_grid(images, title, image_shape=(28, 28), n_col=5, n_row=2):
#     fig, axes = plt.subplots(nrows=n_row, ncols=n_col, figsize=(2. * n_col, 2.
↪26 * n_row))
#     axes = axes.flatten() # Flattening the axes array to simplify the
↪indexing
#     for i, comp in enumerate(images):
#         ax = axes[i]
#         ax.imshow(comp.reshape(image_shape), cmap=plt.cm.gray_r,
↪interpolation='nearest')
#         ax.set_xticks(())
#         ax.set_yticks(())

#     fig.suptitle(title)
#     plt.show()


# Re-running the NMF and plotting for r in {3, 6, 10}
for r in [3, 6, 10]:
    apply_nmf_and_plot(data, r)

```

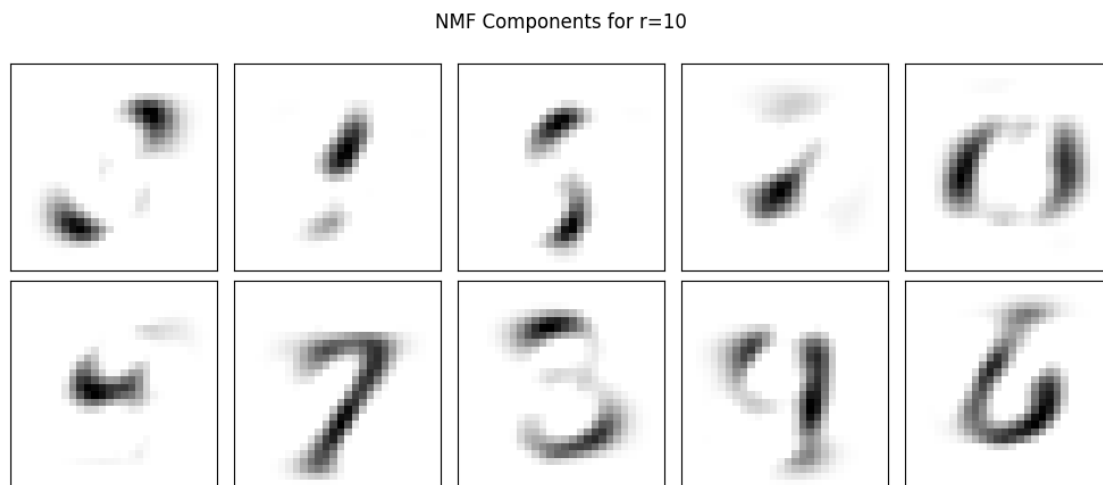


/home/karanvora/miniconda3/lib/python3.8/site-

```
packages/sklearn/decomposition/_nmf.py:1710: ConvergenceWarning: Maximum number
of iterations 200 reached. Increase it to improve convergence.
warnings.warn(
```



```
/home/karanvora/miniconda3/lib/python3.8/site-
packages/sklearn/decomposition/_nmf.py:1710: ConvergenceWarning: Maximum number
of iterations 200 reached. Increase it to improve convergence.
warnings.warn(
```



```
[24]: # Correcting the dimension mismatch error in the compute_nearest_neighbors_NMF
      ↪ function
      def compute_nearest_neighbors_NMF(train, testImage, H):
```

```

    projected_train = [np.dot(np.array(i, dtype=float), H.T) for i in train] #
    ↪Corrected this line
    projected_test = np.dot(np.array(testImage, dtype=float), H.T) # Corrected
    ↪this line
    digit = 0
    imageIdx = 0
    dist = np.linalg.norm(projected_train[digit][imageIdx] - projected_test)

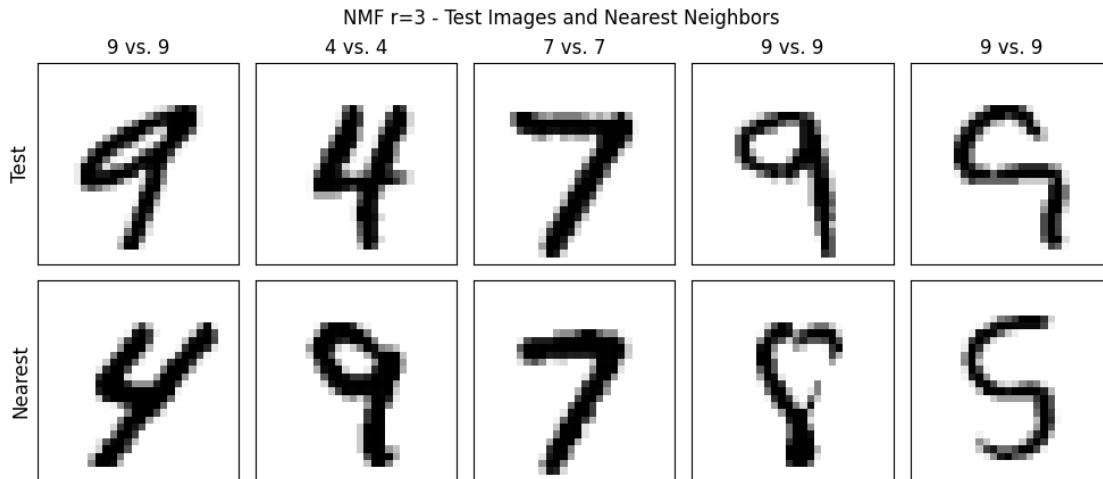
    for i in range(len(train)):
        for j in range(train[i].shape[0]):
            tempDist = np.linalg.norm(projected_train[i][j] - projected_test)
            if tempDist < dist:
                digit = i
                imageIdx = j
                dist = tempDist
    return digit, imageIdx

# Apply NMF with r=3
r = 3
nmf = NMF(n_components=r, init='random', random_state=0)
W = nmf.fit_transform(trainarr)
H = nmf.components_

# Find the nearest neighbors for all test images again
nmf_imgs = []
for i in range(len(testLabels)):
    nnDig, nnIdx = compute_nearest_neighbors_NMF(train, test[i], H)
    nmf_imgs.extend([test[i], train[nnDig][nnIdx,:]])

# Visualize the test images and their nearest neighbors
col_titles = ['%d vs. %d'%(i,j) for i,j in zip(testLabels,testLabels)]
plot_image_grid(nmf_imgs, "NMF r=3 - Test Images and Nearest Neighbors",
                (28,28), len(testLabels), 2, True, row_titles=['Test',
    ↪'Nearest'], col_titles=col_titles)

```



```
[22]: # Correcting the dimension mismatch error in the compute_nearest_neighbors_NMF
      ↪function
def compute_nearest_neighbors_NMF(train, testImage, H):
    projected_train = [np.dot(np.array(i, dtype=float), H.T) for i in train] #
    ↪Corrected this line
    projected_test = np.dot(np.array(testImage, dtype=float), H.T) # Corrected
    ↪this line
    digit = 0
    imageIdx = 0
    dist = np.linalg.norm(projected_train[digit][imageIdx] - projected_test)

    for i in range(len(train)):
        for j in range(train[i].shape[0]):
            tempDist = np.linalg.norm(projected_train[i][j] - projected_test)
            if tempDist < dist:
                digit = i
                imageIdx = j
                dist = tempDist
    return digit, imageIdx

# Apply NMF with r=3
r = 6
nmf = NMF(n_components=r, init='random', random_state=0)
W = nmf.fit_transform(trainarr)
H = nmf.components_

# Find the nearest neighbors for all test images again
nmf_imgs = []
for i in range(len(testLabels)):
    nnDig, nnIdx = compute_nearest_neighbors_NMF(train, test[i], H)
```




```

nmf_imgs.extend([test[i], train[nnDig][nnIdx,:]])

# Visualize the test images and their nearest neighbors
col_titles = ['%d vs. %d'%(i,j) for i,j in zip(testLabels,testLabels)]
plot_image_grid(nmf_imgs, "NMF r=6 - Test Images and Nearest Neighbors",
                (28,28), len(testLabels), 2, True, row_titles=['Test',
↪ 'Nearest'], col_titles=col_titles)

```

NMF r=3 - Test Images and Nearest Neighbors



```

[25]: # Correcting the dimension mismatch error in the compute_nearest_neighbors_NMF
↪function
def compute_nearest_neighbors_NMF(train, testImage, H):
    projected_train = [np.dot(np.array(i, dtype=float), H.T) for i in train] #
↪Corrected this line
    projected_test = np.dot(np.array(testImage, dtype=float), H.T) # Corrected
↪this line
    digit = 0
    imageIdx = 0
    dist = np.linalg.norm(projected_train[digit][imageIdx] - projected_test)

    for i in range(len(train)):
        for j in range(train[i].shape[0]):
            tempDist = np.linalg.norm(projected_train[i][j] - projected_test)
            if tempDist < dist:
                digit = i
                imageIdx = j
                dist = tempDist
    return digit, imageIdx

# Apply NMF with r=3
r = 10

```

```

nmf = NMF(n_components=r, init='random', random_state=0)
W = nmf.fit_transform(trainarr)
H = nmf.components_


# Find the nearest neighbors for all test images again
nmf_imgs = []
for i in range(len(testLabels)):
    nnDig, nnIdx = compute_nearest_neighbors_NMF(train, test[i], H)
    nmf_imgs.extend([test[i], train[nnDig][nnIdx,:]])

# Visualize the test images and their nearest neighbors
plot_image_grid(nmf_imgs, "NMF r=10 - Test Images and Nearest Neighbors",
                (28,28), len(testLabels), 2, True, row_titles=['Test', 'Nearest'], col_titles=col_titles)

```

/home/karanvora/miniconda3/lib/python3.8/site-packages/sklearn/decomposition/_nmf.py:1710: ConvergenceWarning: Maximum number of iterations 200 reached. Increase it to improve convergence.

warnings.warn(



For C, It is clear with PCA the matches are far more closer and accurate than NMF. NMF requires higher r number to get accurate enough Information.